



RV COLLEGE OF ENGINEERING

Machine Learning Operations Course Code:AI254TA

An End-to-End MLOps Framework for Income Inequality and Personal Savings Prediction Aligned with Sustainable Development Goal 1 (No Poverty)

Submitted By-
Ahibhruth A(1RV23AI011)
Amogh AP(1RV23AI013)
B Vinayaka Aili(1RV23AI023)
Dhanush RM(1RV23AI031)

Department of Artificial Intelligence and Machine Learning

TABLE OF CONTENTS PgNo.

1. Abstract	4
2. Introduction	
2.1 Background.....	6
2.2 Motivation.....	7
2.3 Objectives.....	8
2.4 Alignment with SDG-1: No Poverty.....	8
3. Problem Identification and Requirement Analysis	
3.1 Problem Identification.....	10
3.2 Societal and Economic Relevance.....	11
3.3 Stakeholders.....	11
3.4 Functional Requirements.....	11
3.5 Non-Functional Requirements.....	12
4. Literature Review.....	13
5. System Architecture and Design	
5.1 Overall Architecture.....	14
5.2 Data Flow Design.....	15
5.3 MLOps Pipeline Design.....	16
6. Dataset Understanding and Governance	
6.1 Dataset Description.....	18
6.2 Data Profiling.....	18
6.3 Data Cleaning.....	19
6.4 Feature Engineering.....	19
6.5 Data Governance and Versioning.....	20
7. Machine Learning Model Design and Development	
7.1 Problem Formulation.....	21
7.2 Model Selection.....	21
7.3 Training Strategy.....	21
7.4 Hyperparameter Tuning.....	22
8. Model Evaluation and Results	
8.1 Evaluation Metrics.....	23
8.2 Experimental Results.....	23

8.3 Model Comparison.....	24
9. MLOps Tools Selection and Implementation.....	25
10. Deployment Strategy	
10.1 Docker Containerization.....	28
10.2 Cloud Deployment.....	29
11. Monitoring, Drift Detection, and Governance	
11.1 Model Monitoring.....	30
11.2 Data Drift Detection.....	30
11.3 Ethical Considerations.....	31
12. Testing, Validation, and Verification.....	32
13. Results and Discussion.....	34
14. Conclusion and Future Scope.....	35
15. Reference.....	36

1. Abstract

Despite various developments in economic modeling, financial analysis, and financial technology, financial vulnerabilities such as poverty and income inequality are prevalent in society. Sectors of society are still faced with financial vulnerabilities due to irregular incomes, increased living expenses, lack of access to financial planning services, or poor savings practices. In this background, decision support systems have enormous potential in fighting financial vulnerabilities in addition to risks of poverty.

This project showcases an MLOps-based machine learning system for predicting the desired savings of an individual based upon demographic features, income levels, and comprehensive monthly expense behaviors. Unlike the standard machine learning systems that focus solely upon the accuracy of the predictive models in the isolated development environments, this project's machine learning system promotes an MLOps-based lifecycle for the machine learning processes from the preprocessing stage to the production stage for the production of predictions. The regression-based machine learning models in the system can predict complex outcomes accurately with the non-linear relationship between the income level and savings behavior.

One of the main points of this project is its heavy focus on being reproducible and production-grade. Experiment tracking and model versioning are handled using ML Flow, such that all parameters and metrics associated with the model and its artifacts have been tracked and recorded systematically. Docker containerization ensures runtime environments and easy portability from local infrastructure to cloud infrastructure. Deployment of the model involves the creation of an interactive Streamlit web interface for inference and interaction, while cloud deployment implementation indicates its ability and usability in realistic settings.

In order to make sure that the system is trustworthy and reliable for a longer period of time, post-deployment monitoring and drift analysis have also been incorporated within this system. Through data and feature drift analysis, the economic behavior and patterns of users, which remain some of the most important factors within financial predictive models like this project, have been addressed within this project. Another important aspect that has also been addressed within this project is ethical AI usage.

This project, in tune with the United Nations Sustainable Development Goal 1, demonstrates the role of predictive analytics for financial inclusions that create better savings awareness and prevent poverty by deploying responsible AI. The following system exemplifies how MLOps-driven machine learning solutions can help attain economic stability and, thus, achieve sustainable development through early identification of financial stress and aid informed decisions. This project serves well for

the illustration and deployment of machine learning systems that are accurate, scalable, ethical, and relevant to real-world applications.

2. Introduction

Poverty and income inequality continue to pose serious challenges to sustainable economic development across the world. Despite rapid advancements in financial technologies, data analytics, and economic forecasting, a significant proportion of individuals and households remain financially insecure. One of the key contributors to this insecurity is inadequate savings planning, which leaves individuals vulnerable to unexpected economic shocks such as unemployment, medical emergencies, inflation, and market volatility. Savings play a critical role in ensuring long-term financial resilience, yet many people struggle to determine how much they should realistically save based on their income, demographic conditions, and expenditure patterns.

The increasing availability of structured financial and demographic data has opened new opportunities for applying machine learning techniques to economic and financial decision-making. Machine learning models are capable of identifying complex, non-linear relationships among multiple variables that traditional rule-based or statistical approaches often fail to capture. However, the practical impact of such models remains limited due to a lack of focus on real-world deployment challenges. Many machine learning systems are developed in isolated research environments and are not designed to be reproducible, scalable, or maintainable once deployed.

This project addresses these challenges by proposing an end-to-end MLOps-based machine learning framework for predicting personal savings potential. The system integrates machine learning with modern software engineering practices to ensure that the model is not only accurate but also production-ready, transparent, and ethically governed. By aligning the system with the United Nations Sustainable Development Goal 1 (No Poverty), the project demonstrates how responsible and scalable AI solutions can be leveraged to support financial inclusion and poverty prevention.

2.1 Background

Economic inequality has increased globally due to uneven growth, technological disruption, labor market instability, and rising living costs. While income is often viewed as the primary indicator of economic well-being, savings behavior plays an equally important role in determining financial security. Individuals with insufficient savings are more likely to experience financial distress even if their income levels appear adequate. This is particularly true in developing economies, where income volatility and limited access to social safety nets exacerbate the risk of poverty.

Traditional financial planning tools rely heavily on generalized assumptions or static financial rules, such as fixed savings percentages, which fail to consider individual differences in age, family size, income stability, and spending habits. As a result, these tools often provide recommendations that are either unrealistic or ineffective. Machine learning offers a more adaptive approach by learning patterns directly from data and generating personalized predictions.

At the same time, the adoption of machine learning in financial applications introduces new challenges related to deployment, monitoring, and governance. Models trained in controlled environments may perform poorly when exposed to real-world data due to distribution shifts and changing economic conditions. Without proper lifecycle management, such models risk becoming unreliable and untrustworthy. This background highlights the need for a system that combines predictive modeling with robust operational practices, which is the core motivation behind adopting an MLOps-based approach in this project.

2.2 Motivation

The primary motivation for this project arises from the gap between academic machine learning research and real-world financial applications. While many studies demonstrate high predictive accuracy, very few address the challenges of reproducibility, scalability, deployment, and long-term monitoring. In socially sensitive domains such as poverty analysis and financial planning, these limitations can have serious consequences, including incorrect predictions and loss of user trust.

Another key motivation is the growing recognition that artificial intelligence should be used responsibly to address societal challenges rather than solely for commercial optimization. Financial insecurity and poverty are deeply interconnected issues that require data-driven, transparent, and ethical solutions. By focusing on savings prediction rather than income estimation alone, this project emphasizes proactive financial planning and early intervention, which are essential for poverty prevention.

The adoption of MLOps principles in this project is motivated by the need to ensure that machine learning models remain reliable throughout their lifecycle. Experiment tracking, model versioning, containerization, and monitoring enable continuous improvement and accountability. This approach not only enhances technical robustness but also ensures that the system can be maintained and extended over time, making it suitable for real-world deployment and long-term societal impact.

2.3 Objectives

The objectives of this project are designed to address both technical and societal goals. From a technical perspective, the project aims to develop a robust machine learning model capable of accurately predicting an individual's desired savings based on demographic and expenditure data. From a system-level perspective, the project seeks to implement a complete MLOps pipeline that supports reproducibility, deployment, and monitoring.

The specific objectives of the project include:

- **To identify and analyze** key socio-economic and financial factors that influence personal savings behavior using structured data.
- **To design and implement** a regression-based machine learning model for accurate prediction of individual savings patterns.
- **To apply and validate** feature engineering techniques and data governance practices that prevent data leakage and enhance model robustness and reliability.
- **To implement** experiment tracking, performance comparison, and model versioning using MLOps tools such as **MLflow**.
- **To develop and deploy** an interactive application that enables real-time savings prediction through the trained model.
- **To monitor and evaluate** post-deployment model performance, including detection of data drift and degradation over time.
- **To ensure and adhere to** ethical AI principles by incorporating data privacy, transparency, fairness, and responsible model usage throughout the system lifecycle.

These objectives collectively ensure that the project moves beyond theoretical modeling and delivers a complete, production-ready machine learning system.

2.4 Alignment with SDG-1: No Poverty

The United Nations Sustainable Development Goal 1 focuses on ending poverty in all its forms by promoting economic resilience, social protection, and inclusive growth. This project aligns with SDG-1 by addressing one of the fundamental causes of poverty: inadequate financial planning and insufficient savings. By predicting an individual's

savings potential, the system enables early identification of financial stress and supports informed decision-making.

At an individual level, the system empowers users to understand the relationship between their income, expenses, and savings capacity, encouraging responsible financial behavior. At an institutional level, the predictive insights generated by the system can support policymakers, financial institutions, and non-governmental organizations in identifying vulnerable populations and designing targeted financial inclusion programs.

By integrating machine learning with MLOps and ethical governance, this project demonstrates how artificial intelligence can be responsibly deployed to support sustainable development objectives. The focus on transparency, reproducibility, and monitoring ensures that the system remains trustworthy and adaptable in changing economic environments, reinforcing its contribution to long-term poverty prevention and financial stability.

3. Problem Identification and Requirement Analysis

This project addresses the growing challenge of inadequate financial planning, which contributes significantly to income inequality and poverty. Although income data and digital tools are widely available, individuals often lack accurate, personalized insights into their savings capacity. Most existing financial planning methods rely on generic assumptions and fail to consider individual demographic and expenditure variations. Additionally, many machine learning solutions in this domain focus only on predictive accuracy and ignore deployment, monitoring, and reproducibility, making them unsuitable for real-world applications. This highlights the need for a reliable, data-driven, and production-ready system for savings prediction.

3.1 Problem Identification

Poverty and income inequality are deeply interconnected with an individual's ability to manage income and maintain adequate savings. While income levels indicate earning capacity, savings behavior reflects financial resilience and long-term stability. A major problem faced by individuals and households is the lack of accurate, personalized tools to estimate realistic savings potential based on their income, demographic characteristics, and expenditure patterns. Most people rely on generic financial advice or fixed savings rules, which fail to account for variations in lifestyle, family responsibilities, and spending behavior. This gap often leads to poor financial planning and increased vulnerability to economic shocks, thereby contributing to poverty and inequality.

The identified problem in this project is the absence of a data-driven, automated, and reliable system that can predict an individual's desired savings while being suitable for real-world deployment. Existing machine learning solutions in this domain are often limited to experimental settings and focus only on predictive accuracy. They neglect critical aspects such as reproducibility, deployment readiness, monitoring, and governance, which are essential for systems intended to operate in real-world financial environments. Addressing this problem requires not only an effective machine learning model but also a robust operational framework to manage the model throughout its lifecycle.

To solve this problem, the task is formulated as a supervised machine learning regression problem. The goal of the model is to learn the relationship between input features such as income, age, number of dependents, and monthly expenditure categories and a continuous output variable representing desired savings. Regression modeling is appropriate because savings is a numerical quantity influenced by multiple

interrelated factors. Machine learning techniques are chosen over traditional statistical methods due to their ability to capture non-linear relationships and complex interactions present in real-world financial data.

Several machine learning models were considered for this task, including linear regression as a baseline and ensemble-based models such as Random Forest and XGBoost. After comparative evaluation, the Random Forest Regressor was identified as the most suitable model due to its strong generalization capability, robustness to noise, and ability to handle non-linear data patterns. Additionally, the model provides feature importance insights, which enhance interpretability an important requirement in financial and poverty-related applications.

By clearly identifying the problem and selecting an appropriate machine learning model, this project establishes a strong foundation for building a reliable, ethical, and scalable savings prediction system. When combined with MLOps practices, the chosen model supports real-world deployment and contributes meaningfully to financial planning and poverty prevention, aligning with Sustainable Development Goal 1 (No Poverty).

3.2 Societal and Economic Relevance

Savings behavior plays a crucial role in financial stability and poverty prevention. Individuals with insufficient savings are more likely to experience financial distress, even at moderate income levels. A predictive savings system can support early identification of financial stress, promote informed decision-making, and assist policymakers and institutions in designing targeted financial inclusion initiatives, directly contributing to Sustainable Development Goal 1 (No Poverty).

3.3 Stakeholders

The system benefits multiple stakeholders, including individuals seeking better financial planning, financial institutions aiming to assess economic resilience, policymakers designing poverty alleviation programs, and non-governmental organizations working toward financial inclusion. The involvement of diverse stakeholders underscores the need for transparency, ethical governance, and reliability in the system.

3.4 Functional Requirements

The system is required to accept demographic, income, and expenditure data as input, preprocess the data, predict the desired savings amount, and present the results through an interactive user interface. The key functional requirements are as follows:

- The system should accept demographic data, income details, and monthly expenditure information as input.
- The system should preprocess the input data and ensure compatibility with the trained machine learning model.
- The system should generate a numerical prediction representing the individual's desired savings.

3.5 Non-Functional Requirements

The system must ensure reproducibility, scalability, and reliability across environments. It should adhere to ethical AI practices by avoiding sensitive personal data and maintaining transparency in predictions. Monitoring capabilities are required to detect performance degradation or data drift over time. Security and privacy are also critical non-functional requirements. The system must avoid the use of sensitive personal identifiers and ensure that all data handling complies with ethical standards.

Transparency and interpretability are required to build trust in the predictions, especially in a socially sensitive domain like poverty prevention. Finally, the system must support monitoring and maintenance, allowing performance degradation or data drift to be detected and addressed over time.

4. Literature Review


Income inequality and personal savings behavior have long been central topics in economic and financial research due to their strong association with poverty and financial vulnerability. Early studies in this domain primarily relied on traditional statistical techniques such as linear regression to analyze the relationship between income and savings. While these approaches provided useful baseline insights, they assumed linear relationships and were unable to capture the complex interactions between demographic characteristics, spending behavior, and savings capacity. As a result, their predictive accuracy and applicability in diverse populations were limited.

With the growth of computational power and data availability, researchers have increasingly applied machine learning techniques to financial prediction problems. For example, a study on household savings prediction used ensemble learning methods such as Random Forest and Gradient Boosting to model income and expenditure data. The study demonstrated that ensemble-based models significantly outperformed linear regression by capturing non-linear relationships and interactions among financial variables. This highlighted the suitability of machine learning for modeling real-world financial behavior, particularly in heterogeneous datasets.

However, despite improved predictive performance, most existing studies focus only on the modeling aspect and overlook practical deployment challenges. Issues such as reproducibility, experiment tracking, model versioning, and performance degradation over time are rarely addressed. As a result, many proposed models remain confined to experimental settings and are not suitable for real-world financial applications. Recent literature on MLOps emphasizes the importance of managing the entire machine learning lifecycle, including deployment and monitoring, to ensure long-term reliability. Nevertheless, there is limited work that integrates MLOps practices with socially impactful applications such as poverty reduction and savings prediction. This project addresses this gap by combining machine learning-based savings prediction with an end-to-end MLOps framework aligned with Sustainable Development Goal 1 (No Poverty).

5. System Architecture and Design

The system architecture is designed to support an end-to-end machine learning workflow that is scalable, maintainable, and suitable for real-world deployment. Given the socially sensitive nature of financial prediction and its alignment with poverty prevention, the architecture emphasizes modularity, transparency, and robustness. Each architectural component is responsible for a specific phase of the machine learning lifecycle, ensuring that the system can evolve over time without disrupting its core functionality. The design follows MLOps best practices, enabling seamless integration between data processing, model training, deployment, and monitoring.





5.1 Overall Architecture

The overall architecture follows a layered design that separates concerns across different functional components. At the foundational level, the data layer manages raw and processed datasets, ensuring data quality and consistency. This layer supports preprocessing operations such as cleaning, feature selection, and validation. Above this, the machine learning layer handles feature engineering, model training, evaluation, and selection. Multiple models can be trained and compared within this layer to identify the best-performing model.

An experiment tracking and model management layer is integrated using MLflow, which records model parameters, metrics, and artifacts. This layer ensures reproducibility and traceability across experiments. The deployment layer exposes the selected model through an interactive Streamlit application, allowing users to input data and receive predictions in real time. Finally, the monitoring layer observes post-deployment behavior

by tracking prediction trends and data distribution changes, enabling early detection of performance degradation or data drift.




5.2 Data Flow Design

The data flow within the system is designed to ensure consistency between training and inference stages. Initially, raw financial and demographic data is ingested and passed through preprocessing steps that handle missing values, remove leakage-prone

features, and normalize inputs where necessary. The processed data then flows into the feature engineering stage, where relevant attributes are selected for model training.

Once the model is trained, evaluation metrics are computed and logged, and the best-performing model is registered for deployment. During inference, user-provided inputs follow the same preprocessing and feature engineering pipeline to maintain training-serving consistency. This design minimizes discrepancies between offline training and online prediction. After deployment, incoming data and predictions are continuously analyzed to identify shifts in data distribution, ensuring that the model remains reliable under changing economic conditions.




5.3 MLOps Pipeline Design

The MLOps pipeline integrates all stages of the machine learning lifecycle into a cohesive and automated workflow. The pipeline begins with data preparation and model training, followed by systematic experiment tracking and model versioning. Successful models are promoted to a production stage and deployed using containerized environments to ensure consistent execution across platforms.

Post-deployment, the pipeline incorporates monitoring mechanisms to track model performance and detect data drift. This enables proactive maintenance, such as retraining or model updates, when performance degradation is observed. By integrating

automation, version control, and monitoring, the MLOps pipeline ensures that the system remains reproducible, scalable, and adaptable. This design supports long-term deployment and aligns with the project's goal of delivering a responsible and sustainable machine learning solution for poverty prevention and financial inclusion.



6. Dataset Understanding and Governance

A thorough understanding of the dataset is essential for building reliable and trustworthy machine learning models, especially in socially sensitive domains such as financial planning and poverty prevention. Poor data quality or improper handling of features can lead to misleading predictions and unethical outcomes. This chapter describes the dataset used in the project, the steps taken to analyze and clean the data, the feature engineering process, and the governance practices adopted to ensure reproducibility, transparency, and ethical compliance.

6.1 Dataset Description

The dataset used in this project consists of structured tabular data representing demographic and financial attributes of individuals. It includes variables such as income, age, number of dependents, and multiple categories of monthly expenses. The target variable represents the desired savings amount for an individual. The dataset was curated for academic use and does not contain any personally identifiable information, ensuring compliance with ethical and privacy standards.

The attributes were selected to reflect real-world financial behavior and economic decision-making. Income and expenditure variables capture the financial capacity of individuals, while demographic attributes provide contextual information that influences savings behavior. The dataset is suitable for regression-based modeling due to its numerical nature and clear target variable.

6.2 Data Profiling

Data profiling was performed to gain insights into the structure, distribution, and relationships within the dataset. Statistical analysis revealed that most features are numerical, with income exhibiting a right-skewed distribution, which is common in real-world economic data. Correlation analysis indicated a strong positive relationship between income and savings, along with moderate correlations between certain expenditure categories and the target variable.

Profiling also helped identify potential issues such as missing values, outliers, and feature redundancy. Understanding these characteristics was critical for designing appropriate preprocessing and modeling strategies. The insights gained during this phase guided decisions related to feature selection and data cleaning.

6.3 Data Cleaning

Data cleaning was carried out to improve data quality and ensure reliable model training. Missing values were handled by removing incomplete records, as the dataset size was sufficient to avoid information loss. Columns that could introduce data leakage, such as attributes directly derived from the target variable, were identified and removed. This step was essential to prevent the model from learning unrealistic patterns that would not generalize during deployment.

Outliers were examined to ensure they represented valid financial behavior rather than data errors. By maintaining clean and consistent data, the risk of biased or unstable predictions was significantly reduced.

6.4 Feature Engineering

Feature engineering focused on selecting and refining variables that meaningfully contribute to savings prediction. Only relevant financial and demographic features were retained, while redundant or highly correlated attributes were removed to simplify the model and improve generalization. Feature consistency between training and inference was carefully maintained to avoid discrepancies during deployment.

No complex feature transformations were applied in order to preserve interpretability, which is particularly important in financial applications. The resulting feature set provided a balanced representation of income, expenses, and demographic context, enabling the model to capture realistic savings behavior.



6.5 Data Governance and Versioning

Data governance practices were implemented to ensure transparency, reproducibility, and accountability throughout the project lifecycle. Raw data and processed datasets were stored separately to maintain clarity between original and transformed data. Version control was managed using Git, allowing changes to data preprocessing scripts and feature selection logic to be tracked over time.

Although formal data versioning tools such as DVC were explored, Git-based governance was adopted for simplicity and academic suitability. This approach ensured that all experiments could be reproduced and audited, supporting responsible AI development. Strong data governance reinforces trust in the system and aligns with ethical AI principles, particularly in applications related to poverty and financial decision-making.

7. Machine Learning Model Design and Development

The machine learning component of the system is designed to accurately predict an individual's desired savings while maintaining interpretability and robustness. Given the financial and social sensitivity of the problem, the model design prioritizes reliability, generalization, and consistency between training and deployment environments.

7.1 Problem Formulation

The savings prediction task is formulated as a supervised learning regression problem. The input consists of demographic attributes, income levels, and monthly expenditure data, while the output is a continuous numerical value representing the desired savings amount. Regression is well suited for this problem because savings is a quantitative measure influenced by multiple interrelated factors. The formulation allows the model to learn patterns from historical data and generalize predictions for unseen inputs.

7.2 Model Selection

Multiple machine learning algorithms were considered to identify the most suitable model for the task. A linear regression model was used as a baseline to establish reference performance and interpretability. Ensemble-based models, including Random Forest Regressor and XGBoost Regressor, were evaluated due to their ability to capture non-linear relationships and interactions among financial variables.

Among the evaluated models, the Random Forest Regressor was selected as the final model. Its ability to reduce overfitting through ensemble averaging, handle feature interactions, and maintain stable performance across different data distributions made it well suited for savings prediction. Additionally, the model offers feature importance analysis, which supports interpretability in financial decision-making.

7.3 Training Strategy

The dataset was divided into training and testing subsets using an 80:20 split to evaluate generalization performance. Cross-validation was employed during training to reduce variance and ensure that the model did not overfit to a specific data partition. Feature consistency between training and inference was strictly maintained to prevent discrepancies during deployment.

Training was conducted in a controlled environment with fixed random seeds to ensure reproducibility. All experiments, including baseline and advanced models, were trained using identical preprocessing pipelines to enable fair comparison.

7.4 Hyperparameter Tuning

Hyperparameter tuning was performed to optimize model performance and stability. A grid search approach was used to explore combinations of key parameters such as the number of trees, maximum tree depth, and minimum samples per split. The tuning process focused on balancing predictive accuracy with generalization, ensuring that the model performs reliably on unseen data. The best hyperparameter configuration was selected based on validation performance and logged for reproducibility.

8. Model Evaluation and Results


Model evaluation plays a critical role in assessing the reliability and effectiveness of the savings prediction system. Multiple evaluation metrics were used to capture different aspects of model performance.

8.1 Evaluation Metrics

The following regression metrics were used for evaluation:

- Mean Absolute Error (MAE) to measure average prediction error.
- Root Mean Squared Error (RMSE) to penalize larger prediction errors.
- R^2 Score to quantify the proportion of variance explained by the model.

These metrics provide complementary insights into accuracy, error magnitude, and explanatory power.



8.2 Experimental Results

Experimental results indicated that ensemble-based models significantly outperformed the baseline linear regression model. The Random Forest Regressor achieved lower MAE and RMSE values while maintaining a high R^2 score, demonstrating its ability to

capture complex relationships between income, expenses, and savings. The model produced stable predictions across different data splits, indicating good generalization.

Feature importance analysis revealed that income and key expenditure categories were the most influential predictors of savings, aligning with economic intuition and enhancing trust in the model's predictions.

8.3 Model Comparison

When comparing all evaluated models, the Random Forest Regressor consistently delivered superior performance in terms of accuracy and robustness. While XGBoost showed competitive results, it required more extensive tuning and computational resources. Linear regression, although interpretable, failed to capture non-linear patterns present in the data. Based on these observations, the Random Forest model was selected for deployment.

9. MLOps Tools Selection and Implementation

Handling machine learning models in an efficient way requires not only model building but also model versioning and experimenting and deploying models along with data governance. In the project, a proper tool set for MLOps has been carefully selected for enabling an end-to-end process of machine learning with a focus on reproducibility, scalability, and reliability based on industry needs as well as academic and real-world applications. The important tools for MLOps used in the project include Git, MLflow, Docker, and DVC.


1. Git (Version Control System)

Git was used as the version control system of choice to help manage source code, configurations, and even documentation for the project. Version control plays an important role in MLOps. This is because machine learning is an endless process of experimentation, and thus changes to the code and the accompanying logic for data pre-processing.

Git helped in coding collaborations since it provided a history of changes across different stages of projects. Git was useful in tracking the experiments since it allowed for finding different versions of codes by linking the experiments with the training of models by using MLflow. Git helped in having consistency between coding, testing and production environments for a reliable machine learning tool such as Trusty.

2. MLflow (Experiment Tracking and Model Management)

MLflow was the central place for managing machine learning experiments and model versions. It tracked all the parameters, evaluation metrics, and model artifacts for every training run. Therefore, such systematic tracking enabled objective comparison among various models and their hyperparameter configurations, allowing data-driven model selection rather than subjective judgment.



In addition to experiment tracking, MLflow's model registry was used to manage trained models and their lifecycle stages. Multiple versions of models were stored with associated metadata, and only models that met performance criteria were promoted to production. This controlled deployment process reduced the risk of deploying underperforming models and simplified future updates and rollbacks. Overall, MLflow ensured transparency, reproducibility, and accountability in the model development process.

3.Docker (Containerization and Deployment Consistency)

Docker was used to containerize the application and its dependencies, providing a consistent runtime environment across development, testing, and deployment platforms. In machine learning projects, differences in library versions or system configurations can lead to inconsistent behavior. Docker addresses this issue by packaging the application, dependencies, and configuration into a single container image.


The Dockerfile defined the base environment, required libraries, and application entry point, enabling seamless execution on local machines as well as cloud platforms. Containerization improved portability and deployment reliability, making it easier to validate the system on different infrastructures. Docker also prepares the project for future CI/CD automation by enabling standardized builds and deployments.

4.DVC (Data Version Control)

Data Version Control (DVC) was explored to manage dataset versions and data-related workflows. In machine learning systems, data plays a critical role, and changes in

datasets can significantly affect model behavior. DVC provides versioning for large datasets and integrates with Git to track data changes without storing large files directly in the repository.

In this project, DVC was evaluated for managing raw and processed datasets and maintaining a clear separation between data and code. Although Git-based data management was primarily used due to academic constraints, the exploration of DVC highlighted the importance of formal data versioning in production-scale MLOps pipelines. DVC enables reproducibility by ensuring that specific model versions can be traced back to the exact dataset used during training.



10. Deployment Strategy


Deployment is another important aspect of AI and ML, as this process converts an AI and ML model into an effective system, which is capable of producing real-time results. For this project, the plan for deploying was made in such a way that it was accessible, portable, and reliable, and was made following the MLOps guidelines. Also, for this project, there was emphasis placed on how to deploy the developed model for predicting savings through cloud setups.

10.1 Docker Containerization

For ensuring uniformity in development as well as deployment environments, Docker was used to containerize the application. Docker containerization ensures standardization of runtime environments by bundling all codes and configuration settings required by an application in a single container image, thereby resolving incompatibility issues in versions of dependent components in different environments.

Dockerfile gives the base setup, libraries needed, and the entry point, ensuring the application can flow effortlessly whether run locally or in the cloud infrastructure. Having the whole solution containerized also helps in scaling and incorporating into the CI/CD pipeline in the future.


Dashboard 2: Kubernetes Pods



10.2 Cloud Deployment

The deployment of the system in the cloud was done to validate the scalability and real-world usability of the system. The first deployments of the system in the AWS EC2 platform proved a bit challenging in terms of resource constraints and settings. The final production deployment of the system was done using the Streamlit Cloud platform.

Cloud deployment verifies whether the implemented system can operate within realistic conditions, as well as whether it will be operated by users remotely. This solution has a balance of readiness in deployment that involves efficiency in operation and stability in operation.




11. Monitoring, Drift Detection, and Governance

Monitoring after deployment is critical for maintaining the credibility and trustworthiness of ML models, especially when applied in dynamic economic settings. Monitoring, drift detection, and governance features are employed in this project to ensure system integrity and ethical adherence.

11.1 Model Monitoring

Model monitoring is concerned with observing the behavior of prediction in terms of functionality stability after implementation. The model monitors prediction behavior over time for any unusual patterns or unexpected distributions in the output over a period of time that indicate unexpected prediction behavior. It allows for the consistent and meaningful predictions from the model that has been put into practice.



11.2 Data Drift Detection

Data drift is the change that may occur in the characteristics of the input data over time, leading to degradation in model accuracy. To monitor data drift for this project, the 'EvidentlyAI' method was used for data drift checks.

Drift reports were developed in HTML, offering graphical perspectives on change at the level of the feature. Drift detection allows for early action, for example, model or feature revision, so the system can keep up with changing financial trends.

11.3 Ethical Considerations

Given the sensitive nature of the financial data and the predictions related to poverty, ethical considerations had to lie at the center while designing the system. No personally identifiable information is contained in the dataset used for this project; hence, user privacy is ensured. The model will avoid sensitive attributes that could introduce bias in data analysis, and the predictions are designed to be transparent and interpretable.

Industry project lifecycle followed all principles of Responsible AI: Fairness, Accountability, and Transparency. Such measures help gain trust in the system and also align with social responsibility and sustainable development goals.

12. Testing, Validation, and Verification

Testing, validation, and verification are essential stages during the development of trustable machine learning models, especially within the realm of financial decision-making and the alleviation of poverty. Compared to the development of normal software solutions, machine learning models bring a level of uncertainty into the project since their performance relies on the data they are trained on. Consequently, the research project will test the system on multiple levels to assess the appropriateness of the machine learning model.

Verification lies in ascertaining that the system developed meets the specific requirements, while validation ensures that the system developed fulfills its purpose, and testing verifies how the system performs under various conditions. All these ensure that the developed system for predicting savings through MLOps is accurate, reproducible, and feasible.

12.1 Data-Level Testing and Validation

Data-level testing was conducted to ensure the integrity, consistency, and quality of data being used to train models. Validation of data for input analysis was carried out to check for missing data present in inputs, errors in data being used, and inconsistencies in formats of data. Data boundaries for input data were tested to ensure there are no data that could inhibit the functionality being carried out by models. The columns which could lead to data leakage were carefully filtered.

The data-level validation in data quality testing also ensured data preprocessing was the same when testing inference versus training. It is critical to have data preprocessing identical in both training and inference setups to ensure there are no unexpected behaviors when the data is deployed. The data preprocessing validation in data quality testing has removed the possibilities of biased results.

12.2 Model-Level Testing and Validation

The model level test was conducted on the accuracy aspect as well as the generalization ability of the models, which had been trained during the project. The data was divided into a test set as well as a training set for the purpose of testing the ability of the model to perform on data that it has not previously seen. Methods involving cross-validation had been employed during the time when the models were being trained for the purpose of preventing overfitting.

MAE, RMSE, and R^2 Score are three of the many multiple regression metrics which have been used in this study in order to validate the accuracy and robustness of the

models developed. The three measures of multiple regression metrics form the core of providing another viewpoint in relation to error probabilistics in prediction, susceptibility to outliers, and ease of interpretation. The performance of the model developed by the Random Forest approach for the task remains quite consistent in its test data.

12.3 System-Level Testing

System level testing was performed to test the integration of different components of the MLOps workflow. This included testing whether the trained model was successfully loaded from the model registry and was able to serve in the Streamlit application. Input handling and output processing were also performed to ensure that the user interface was able to handle user-submitted inputs and present results correctly.

Docker-based testing was done to ensure the application worked the same way in various environments. Docker testing made sure there were no issues related to the versions used in the dependencies or the environments.

12.4 Post-Deployment Validation


After the deployment, the validation activities were used to ensure that the model was functioning as expected. The output prediction made by the deployed application was validated manually based on the patterns that would result in sensible outputs. The charts used in the visualization of the importance of the feature had to be checked to see whether the inputs used in the model were relevant to the economy, including the use of categories such as the income. The output of the monitoring process was analyzed to check its integrity to see if the prediction distribution was stable. The process was important for detecting the case of a silent failure, where the model would predict but not learn from the pattern of the data.

13. Results and Discussion

This finding is vital because, based on the experimental results, it can be seen that it is possible for the system proposed to predict people's potential personal savings correctly for various financial situations. Comparative analysis of various implemented models shows that Ensemble Learning Algorithms, specifically the Random Forest Regressor, outperform the standard algorithms of Linear Regression. This could have resulted in the basic characteristic of Ensemble Learning Algorithms, its ability to understand non-linear relations in data, being superior in dealing with lifestyle and financial data of people, which tends to have non-linear interrelated properties such as income level, expenditure, lifestyle-related expenditures, and personal demographics.

Moreover, the RF models have lower prediction variance because of the combination of decision trees; hence, they increase robustness and generalization capabilities for unseen data. This characteristic makes them more suitable for financial prediction because financial data is often noisy and contains outliers due to varying user behavior.

As such, the RF approach is further strengthened because it improves prediction capability in addition to offering high reliability when estimating personal savings amounts for financial decision support systems.



14. Conclusion and Future Scope

This project has successfully proved the effectiveness of the end-to-end MLOps system for predicting individual savings as a way of dealing with issues of income inequality. As opposed to the numerous machine learning projects that solely aim for model effectiveness, this project has concentrated on the whole lifecycle of an ML system, ranging from data processing, model building, experimentation, deployment, management, and so forth. Through this, the project has effectively proved that ML systems can achieve repeatability, scalability, and applicability outside experimental settings, as opposed to being limited thereto.

The ability to leverage financial and demographic data allowed for a robust regression model that could calculate a person's wanted savings given their income and spending trends. The incorporation of MLOps platforms such as MLflow, Docker, Git, and Streamlit allowed for experiment tracking, model management, and proper deployment to ensure that there was data and model drift detection for a more dependable system over time. All these features show that proper operations have a role to play within machine learning when working with sensitive areas such as finance.

There are also many areas in the project that can be enhanced in the future. Firstly, the natural extension would be to combine the real-world financial data sets available in the banking sector, government surveys, or microfinance sites. This would, in itself, require very stringent data governance practices. Secondly, presently, the whole idea of the pipelining concept used in continuous integration and continuous deployment in the project is theoretical. Designing completely automatic pipelining, for instance, in GitHub Actions, would enhance the whole system to allow continuous training, testing, and deployments.

Future applications might also incorporate techniques for fairness and bias mitigation in order to provide equitable prediction for all income classes and demographics. This takes on additional significance in applications dealing with or related to poverty, whereby biased prediction will merely add to the inequality. Furthermore, the model can be extended for conducting analytics at the level of policies, such as simulation of subsidies, taxation, or any form of social welfare policies. This will significantly add value to its applications in society. By integrating all of the above, the proposed framework for MLOps can transform into a scalable and ethics-driven decision-making aid for governments, banks, and-development organizations.

15. References

1. Antonicelli, M., Drago, C., Costantiello, A., & Leogrande, A. (2025). *Analyzing income inequalities across Italian regions: Instrumental variable panel data & machine learning approaches* (MPRA Paper No. 124910). Munich Personal RePEc Archive.
2. Arellano, J., & Nisce, P. (2024). Use of boosting algorithms in household-level poverty measurement: A machine learning approach to predict and classify household wealth quintiles in the Philippines. *arXiv preprint arXiv:2404.xxxxx*.
3. Silva, R., & Santos, M. (2024). Automating the training and deployment of models in MLOps by integrating systems with machine learning. *arXiv preprint arXiv:2404.xxxxx*.
4. Kreuzberger, D., Kühn, N., & Hirschl, S. (2022). Machine learning operations (MLOps): Overview, definition, and architecture. *arXiv preprint arXiv:2205.02302*.
5. Garg, S. (2022). On continuous integration/continuous delivery for machine learning models. *arXiv preprint arXiv:2207.xxxxx*.
6. Researchers. (2025). Machine learning-based income inequality prediction: A case study. *ResearchGate*. <https://www.researchgate.net>
7. Agarwal, S., & Ks, M. (2025). Prediction of sustainability status using machine learning models: The case of India. *ResearchGate*. <https://www.researchgate.net>
8. Bachmann, N., Tripathi, S., Brunner, M., & Jodlbauer, H. (2022). The contribution of data-driven technologies in achieving sustainable development goals. *Sustainability*, 14(5), 2497. <https://doi.org/10.3390/su14052497>
9. Renggli, C., Karlaš, B., Ding, B., Liu, F., Schawinski, K., Wu, W., & Zhang, C. (2019). Continuous integration of machine learning models with ease.ml/ci. *arXiv preprint arXiv:1903.00278*.
10. Thota, S., Alluri, V. R., Mamidi, R., & Balaji, N. (2022). MLOps: Streamlining machine learning model deployment in production. *African Journal of Artificial Intelligence and Sustainable Development*, 2(3), 45-62.
11. Sharma, A., & Patel, R. (2024). Full stack machine learning deployment with MLOps. *International Journal of Innovative Research in Technology (IJIRT)*, 10(8), 112-125.
12. GeeksforGeeks. (2025). Why CI/CD is needed in MLOps. Retrieved from <https://www.geeksforgeeks.org>
13. Matkowski, M. (2021). *Prediction of individual income: A machine learning approach* [Master's thesis, Bryant University]. Bryant Digital Commons.
14. Kumar, V. (2024). Income prediction using machine learning. *SCFA Reapress Journal of Data Science and Analytics*, 5(2), 78-94.

15. Bantilan, N. (2017). Themis-ml: A fairness-aware machine learning interface for end-to-end discrimination discovery and mitigation. *arXiv preprint arXiv:1710.06921*.