



RV College of Engineering®

Mysore Road, RV Vidyaniketan Post, Bengaluru - 560059, Karnataka, India

Intelligent Workforce Allocation System

MLOps Tutorial Report

submitted by

Ramita K A	1RV23AI079
Rushil Shodavaram	1RV22AI082
Raghavi U Baler	1RV23AI078
Preetham R	1RV23AI075

Artificial Intelligence and Machine Learning

2024-2025

TABLE OF CONTENTS		
Phase 1:Model Development and Foundations		
Chapter 1		
1.1	Problem Statement and Objectives	4
1.2	Project Scope and Evaluation Metrics	5
1.3	Tools and Environmental Setup	6
1.4	Literature Review	7-9
Chapter 2		
2.1	Understanding MLOps	10
2.2	Core MLOps Principles	10
2.3	Requirement Analysis	11
2.4	Risk Matrix	12
2.5	Risk Analysis	12-13
Chapter 3		
3.1	Data Collection and Exploration Techniques	14
3.2	Data Profiling, cleaning and preprocessing	15
3.3	Feature Extraction and Feature Selection	15
3.4	Model Design Approaches	16
3.5	Governance and Data Version Control	16
Chapter 4		
4.1	Model Building	17
4.2	Training, Validation and Testing	17
4.3	Evaluation Metrics and Visualisation	17
4.4	Experiment Tracking	18
4.5	Interpreting Model Results and Error Analysis	19
Phase 2:Deployment, Monitoring and Governance		
Chapter 5		
5.1	Revisiting Model Performance	20
5.2	Hyperparameter Tuning and Performance	20
5.3	Interpretability	21
5.4	Model Versioning and Registry	22
5.5	Documentation of Model Updates	22
Chapter 6		
6.1	Overview of CI/CD for ML	23
6.2	Tools Setup:GitHub Actions	23

6.3	Automating Model Training and Testing	24
6.4	Building and Running a Deployment Pipeline	25
6.5	Continuous Integration Hands-On Exercise	25
Chapter 7		
7.1	Model Monitoring Concepts and KPI's	26
7.2	Detecting Concept Drift and Data Drift	26
7.4	Model Retraining and Redeployment	27
7.5	Logging,Alerting and Dashboard Setup	28
Chapter 8		
8.1	Post-Deployment Performance Review	29
8.2	Continuous Improvement via Feedback Loops	30
8.3	Governance and Ethical AI Practices	30
8.4	Documentation: Model Cards, Data Cards	31
8.5	Audit and Review Checklist	31
Appendix		
A.1	List Of MLOps Tools, Its Features, Pros And Cons	32-34
A.2	Tool Installation Guides	34
A.3	References	35-36
A.4	Folder Structure	36
A.5	Further Reading	37

Chapter 1

Introduction

1.1 Problem Statement and Objectives

1.1.1. Problem Statement

In modern organizational workflows, task allocation is often performed manually or through simplistic rule-based systems that fail to account for the dynamic and multi-dimensional nature of employee capabilities. This leads to:

- **Inefficient resource utilization** – tasks assigned to employees with mismatched skills
- **Burnout and uneven workload distribution** – some employees overburdened while others underutilized
- **Reduced productivity and quality** – tasks not matched to expertise result in delays and errors
- **Lack of transparency and fairness** – allocation decisions appear arbitrary or biased

Traditional workforce management systems lack the intelligence to dynamically match employees to tasks based on real-time availability, historical performance, skill alignment, and task complexity

1.1.2. Objectives

The primary objective of this project is to implement a **Workforce Allocation Allocation System**. This system helps to fundamentally revolutionize the traditional, often manual, process of task assignment within an organization. The core function is to employ advanced analytical and machine learning techniques to automatically assign the **most suitable employee** to any given task.

1. Skill match between employee and task requirements
2. Employee availability and workload
3. Historical performance metrics
4. Task complexity

1.2 Project Scope and Evaluation Metrics

1.2.1. Project Scope

This project focuses on developing an **Intelligent Workforce Allocation System** that predicts employee–task suitability scores using historical data and engineered features. The scope

includes building and evaluating multiple regression models to identify the most effective approach based on performance and interpretability.

The project incorporates **MLOps practices** such as dataset versioning using **DVC**, experiment tracking and reproducibility using **MLflow**, and automated training and monitoring pipelines. Model performance and data drift are continuously assessed using tools like **EvidentlyAI**. The final model is deployed as a **Streamlit-based web application**, containerized with **Docker** to ensure portability, consistency, and real-time inference.

1.2.2. Evaluation Metric

To comprehensively assess both predictive accuracy and operational reliability of the workforce allocation model, distinct metric categories are employed:

1. Regression Performance Metrics:

- **Mean Absolute Error (MAE)**
Quantifies the average absolute deviation between predicted and actual suitability scores. Lower MAE values indicate superior predictive accuracy across the dataset.
- **Root Mean Squared Error (RMSE)**
Computes the square root of the mean squared prediction errors. RMSE penalizes larger prediction deviations more severely than MAE, which is critical when substantial allocation errors could significantly impact workforce efficiency.
- **R² Score (Coefficient of Determination)**
Measures the proportion of variance in suitability scores explained by the model. Higher R² values (approaching 1.0) indicate stronger model performance and better capture of underlying data patterns.
- **Mean Absolute Percentage Error (MAPE)**
Expresses prediction errors as percentages of actual values, facilitating intuitive interpretation and comparison across different scales

2. Model Monitoring and Drift Detection Metrics

- **Data Distribution Drift Measures**
Statistical tests (e.g., Kolmogorov-Smirnov test, Population Stability Index) detect shifts in feature distributions between training and production data, signaling potential degradation in model accuracy.
- **Prediction Drift Analysis**
Monitoring changes in the distribution of model predictions over time helps identify concept drift where the relationship between features and target variable evolves.
- **Performance Stability Monitoring**
Continuous tracking of MAE, RMSE, and R² on incoming production data

enables early detection of performance degradation, triggering retraining protocols when necessary.

3. System Performance and Operational Metrics

- **API Response Latency**
Target: < 200ms for 95th percentile requests, ensuring real-time responsiveness for allocation queries.
- **Throughput Capacity**
The system must handle ≥ 100 concurrent requests per second without performance degradation.
- **Model Retraining Frequency**
Scheduled monthly retraining cycles with drift-triggered ad-hoc retraining when distribution shifts exceed defined thresholds.

1.3 Tools

MLOps stage	Tools	Purpose
Data Management and Versioning	DVC	Track Dataset changes, versions, reproducibility
Experiment Tracking and Model Management	ML Flow	Log experiments,compare runs,store best models
Pipeline Orchestration and Workflow Automation	Prefect	Automate end-to-end ML workflow
Model Deployment and serving	Streamlit	Deploy the trained model as a Streamlit-based web application.
Monitoring,Drift Detection and Governance	Evidently AI	Track data drift , model performance and degradation

1.4 Environmental Setup

The project environment is designed to ensure **reproducibility, scalability, and ease of deployment**. Development is carried out using **Python 3.10** within a virtual environment, with dependencies managed through a centralized **requirements.txt** file. **MLflow** is used for experiment tracking, enabling systematic logging of model parameters, metrics, and artifacts. **DVC** handles dataset versioning to ensure reproducibility while keeping large files out of the Git repository. The ML workflow is orchestrated using **Prefect**, structuring tasks for data loading, training, evaluation, and monitoring. **Evidently AI** is employed to monitor data quality and detect data drift by comparing production and baseline datasets. The trained model is deployed

through a **Streamlit-based web application**, containerized using **Docker** to ensure consistent runtime environments. Finally, **GitHub Actions** supports CI/CD by automating validation and pipeline execution on repository updates.

1.5 Literature Review

A. Sharma and P. Kumar [1], This paper proposes a fully integrated MLOps pipeline using DVC for data versioning and MLflow for experiment tracking. It demonstrates how combining these tools ensures full reproducibility by linking specific Git commits to exact data snapshots (via DVC) and model metrics (via MLflow).

K. Patel and S. J. Rao [2], This paper focuses specifically on DVC (Data Version Control) as a cornerstone for Continuous Training (CT) pipelines. It analyzes how DVC's storage-agnostic architecture allows teams to handle large-scale datasets (images/video) in a Git-like manner, which is essential for maintaining audit trails in regulated industries.

J. Nilsson et al. [3], This thesis implements a fraud detection system that uses Evidently AI for micro-batch monitoring. The study verifies that Evidently's drift detection algorithms can identify "concept drift" in financial transaction data significantly earlier than traditional batch monitoring methods, reducing fraud losses.

R. Müller et al. [4], This study evaluates Evidently AI alongside other drift detection tools using real-world datasets from smart buildings. It concludes that Evidently AI excels in general data drift detection due to its comprehensive statistical tests (like Kolmogorov-Smirnov) and user-friendly visualization reports, making it highly effective for production monitoring.

D. Al-Fraihat et al. [5], To address task allocation challenges in Distributed Agile Software Development, this study evaluated four classifiers: Random Forest, K-Nearest Neighbors (K-NN), Decision Tree, and AdaBoost. Random Forest achieved the highest accuracy at 96.7%, outperforming K-NN (94.2%), Decision Tree (93.5%), and AdaBoost (93%). These results demonstrate that machine learning is a highly effective tool for optimizing role assignments and reducing project risks.

N. Grigoryan [6], This research provides a comparative analysis of orchestration tools including Prefect, Kubeflow, and Apache Airflow. It highlights Prefect's "negative engineering" philosophy (handling failures and retries automatically) and its superior Python-native developer experience compared to the rigid DAG structures of older orchestrators.

S. Sharma and R. Gupta [7], This paper describes a system that uses AI and Machine Learning to offer dynamic prioritization and contextually informed scheduling. It moves beyond traditional manual checklists by adapting to real-time changes in work patterns and deadlines to suggest the most efficient task execution order.

V. Nakra et al. [8], This paper proposes an integrated AI system combining NLP, Reinforcement Learning, and Bayesian Networks to replace manual task allocation. The system extracts task requirements from documents and predicts optimal resource allocation, achieving a 21% improvement in resource utilization and 31% faster project delivery compared to baseline methods.

S. Gupta and R. S. Rao [9], This paper explores the Large-Scale Scrum (LeSS) framework and proposes a predictive model using Multi-Layer Perceptrons and Linear Regression to compute optimal Human Resource (HR) scores. It aims to mitigate challenges in team composition by predicting the best fit for specific security and network engineering tasks.

A. K. S. and P. R. [10], This project develops a web-based application for allocating daily wage workers to tasks using a system dynamics approach. The study concludes that such a system improves communication, optimizes resource usage, and enhances overall productivity by accurately matching employee capabilities to task demands.

H. Grover and S. Sharma [11], This paper proposes an integrated AI system combining NLP, Reinforcement Learning, and Bayesian Networks to replace manual task allocation methods. Testing demonstrated significant performance gains, achieving 21% better resource utilization, 31% faster project delivery, and 83% alignment with expert decisions compared to traditional baselines.

V. T. P. N. and K. C. [12], This research introduces a multi-objective evolutionary algorithm to handle dynamic events, such as employee turnover, in software project scheduling. The model re-optimizes task assignments by considering the skills acquired by employees versus those required by tasks, minimizing the negative impact of staff changes on project cost and duration.

R. J. Wave [13], This paper details a workforce scheduling system that uses a Random Forest classifier to assign employees to projects with 89% accuracy. It also integrates a computer vision-based attendance system using Raspberry Pi to ensure real-time tracking and reduce administrative overhead.

A. Sadegheih et al. [14], This study presents a metaheuristic model that applies fuzzy ranking and genetic algorithms to project scheduling under uncertainty. It models activity durations as fuzzy numbers and uses a parallel fuzzy prioritization method to generate initial populations that optimize resource leveling.

D. Mourtzis et al. [15], The authors present a decision-making framework for allocating multi-skilled operators to specific tasks in industrial environments. The model optimizes the correlation between operator skills and job requirements while minimizing workforce costs, demonstrating effectiveness in a real-life industrial scenario.

P. K. and S. R. [16], This study proposes a Convolutional Neural Network (CNN) model combined with Word2Vec embeddings to automate the "bug triage" process. The system analyzes bug reports and recommends the top 10 most suitable developers to fix the issue, reducing the manual labor involved in assigning bugs.

S. K. and R. B. [17], This review paper analyzes the application of Reinforcement Learning (RL) techniques, such as Q-Learning and Deep Q-Networks, for dynamic task scheduling. It highlights how RL agents can learn optimal policies for resource allocation in uncertain environments where task arrival times and resource availability fluctuate.

M. A. and S. H. [18], This paper introduces a scheduling system based on fuzzy logic that considers input parameters like cost, timeline, and efficiency. The proposed method utilizes triangular membership functions to prioritize tasks and allocate resources, demonstrating better waiting and turnaround times compared to standard algorithms like Min-Min.

M. Zaharia et al. [19], This seminal paper introduces MLflow, an open-source platform designed to manage the end-to-end machine learning lifecycle. It details the architecture for tracking experiments, packaging code for reproducible runs, and sharing models, addressing the "spaghetti code" problem common in ML development.

F. Roque et al. [20], This study formalizes the problem of task allocation in distributed crowdsourcing environments as an AI planning problem. It evaluates whether automated planning can effectively match the heterogeneous skills of crowd workers to specific software engineering tasks to ensure high-quality results.

Recent research emphasizes combining robust MLOps practices with advanced AI to optimize software resource allocation. Studies on tools like MLflow, DVC, Prefect, and Evidently AI highlight the necessity of reproducible, automated pipelines for modern systems. Concurrently, applying machine learning models—such as Random Forest and Reinforcement Learning—to workforce planning has proven to significantly outperform traditional manual methods. These findings demonstrate that integrating predictive analytics not only enhances resource utilization and delivery speed but also aligns closer to expert decision-making than legacy approaches.

This project develops an Intelligent Workforce Allocation System designed to optimize task assignment by predicting an employee's suitability score based on the alignment between their current skills, the required task skills, and their availability. To ensure robust performance, the predictive engine was built by evaluating and comparing Random Forest, XGBoost, and Linear Regression models. This development lifecycle is supported by a comprehensive MLOps pipeline: DVC is utilized for data versioning and reproducibility, MLflow tracks experiments and model metrics, Prefect orchestrates the automated workflow, and Evidently AI monitors the deployed models for data drift, ensuring the system remains reliable and effective in dynamic operational environments.

Chapter 2

Understanding MLOps, Risks, And Requirement Analysis

2.1 What is MLOps?

MLOps is a set of practices that combines Machine Learning, DevOps, and Data Engineering to deploy, monitor, and maintain machine learning models in production environments. MLOps bridges the gap between model development (performed by data scientists) and model deployment (handled by engineers and operations teams), ensuring that ML systems are reliable, scalable, and maintainable throughout their lifecycle.

In the context of the Intelligent Workforce Allocation System, MLOps enables us to:

- Track and version our employee-task datasets as they evolve
- Experiment with different regression models systematically
- Automate the training and deployment pipeline
- Monitor model performance and detect when retraining is needed
- Ensure the system remains fair, accurate, and compliant over time

2.2 Core MLOps Principles

1. Data Management and Versioning

Datasets (profiles, tasks, allocation records) are versioned using **DVC** to ensure reproducibility, track changes, and facilitate conflict-free collaboration.

2. Experiment Tracking and Model Management

MLflow tracks iterative experiments, logging run parameters (e.g., learning rate), metrics (R^2 , MAE, RMSE), storing model artifacts, and tracking the deployed production version.

3. Pipeline Orchestration and Workflow Automation

Prefect automates the end-to-end ML workflow, automatically triggering retraining on new data, performing data validation, orchestrating the sequence.

4. Model Deployment and Serving

Streamlit deploys the workforce allocation model through an interactive web interface, allowing users to input employee and task details and receive real-time suitability scores.

5. Monitoring, Drift Detection, and Governance

Evidently AI generates weekly drift detection reports, monitors performance and prediction distributions, signals retraining needs due to drift, and ensures fair allocation decisions.

2.3 Requirement Analysis

Functional Requirements

ID	Requirement	Description
FR1	Suitability Score Prediction	System accepts employee_id and task_id, returns suitability score (0-1)
FR2	Prediction Endpoint	It accepts user inputs via form fields and returns real-time suitability scores
FR3	Top-N Employee Ranking	Given task_id, return ranked list of N most suitable employees
FR4	Explainable Predictions	Each prediction includes SHAP values showing feature contributions
FR5	Batch Prediction	Support predict_batch for multiple employee-task pairs simultaneously
FR6	Model Versioning	MLflow maintains history of all trained models with rollback capability
FR7	Data Validation	Prefect validates input schemas before training/prediction
FR8	Logging	Log all predictions, response times, and errors for monitoring

Non-Functional Requirements

ID	Requirement	Target	Tool
NFR1	Response Time	< 200ms per prediction	Streamlit
NFR2	System Availability	$\geq 99.5\%$ uptime	Docker
NFR3	Model Accuracy	$R^2 \geq 0.85$, $MAE \leq 0.10$	MLflow
NFR4	Scalability	Handle 100+ concurrent requests	Load Balancing
NFR5	Security	HTTPS + Authentication	Streamlit

2.4 Risk Matrix

Likelihood ↓ / Consequence →	Negligible	Minor	Moderate	Severe
High	—	Model Drift	Model Bias	Data Quality
Medium	Integration Failure	—	Algorithmic Bias	Data Privacy
Low	—	CI/CD Deployment Failure	Scalability Issues	Data Availability
Very Low	Pipeline Failures	Cost Overruns	Human Trust	Legal Non-Compliance

Figure 2.1 : Risk Matrix

A **Risk Matrix** evaluates project risks based on:

- Likelihood (Y-axis): Probability of occurrence
- Impact (X-axis): Severity of consequences

Risks in the top-right (high likelihood + high impact) require immediate action, while bottom-left risks need only monitoring

2.5 Risk Analysis

Critical Risks (High Likelihood + High Impact)

R1: Model Bias

AI assigns tasks unfairly based on historical patterns (e.g., always giving seniors complex tasks).
Mitigation: Monitor with Evidently AI; conduct bias audits; balance training data.

R2: Data Quality

Incomplete employee profiles (missing skills/availability) cause poor predictions.
Mitigation: Validate data in Prefect pipelines; enforce schema checks.

High Risks

R3: Model Drift (High Likelihood + Medium Impact)

Employee skills change over time; the model becomes outdated.
Mitigation: Weekly drift monitoring (Evidently AI); monthly retraining (Prefect).

R4: Algorithm Bias (Medium Likelihood + High Impact)

The model systematically favors certain employee groups.

Mitigation: Test multiple algorithms in MLflow; validate fairness metrics.

R5: Data Privacy (Medium Likelihood + High Impact)

Sensitive employee data may be exposed through the Streamlit web interface.

Mitigation: Implement access control, restrict user roles.

R6: Scalability Issues (Low Likelihood + High Impact)

The system cannot handle the growing workforce.

Mitigation: Horizontal scaling with Docker; load testing.

Medium Risks**R7: Integration Failure (Medium Likelihood + Low Impact)**

MLOps tools fail to integrate properly.

Mitigation: Integration testing; maintain documentation.

R8: CI/CD Deployment Failure (Low Likelihood + Medium Impact)

Automated pipeline prevents model updates.

Mitigation: Prefect retry logic; manual deployment backup.

Low Risks**R9: Pipeline Failures (Very Low Likelihood + Low Impact)**

Mitigation: Automatic retries; alerting.

R10: Cost Overruns (Very Low Likelihood + Medium Impact)

Mitigation: Budget monitoring.

R11: Human Trust (Very Low Likelihood + High Impact)

Mitigation: SHAP explanations; transparency.

R12: Legal Non-Compliance (Very Low Likelihood + High Impact)

Mitigation: Legal review; audit logs.

Chapter -3

Data understanding, Feature engineering and governance

3.1 Data Collection and Exploration Techniques

Data collection is a critical step in building the Intelligent Workforce Allocation System, as the quality of input data directly influences model accuracy and fairness. The dataset used in this project represents historical workforce and task allocation records collected from organizational task management systems. **The dataset primarily consists of:**

- Employee profile data (skills, experience, availability)
- Task-related data (task type, complexity, required skills)
- Historical performance indicators (task completion success, feedback scores)

Data exploration was performed to understand the structure, size, and characteristics of the dataset. Exploratory Data Analysis (EDA) techniques were applied to:

- Identify feature distributions
- Detect missing or inconsistent values
- Analyze relationships between employee attributes and task suitability scores

Statistical summaries and visual inspection techniques helped in identifying skewed features, outliers, and correlations that guided subsequent preprocessing and feature engineering steps.




Figure 3.1: Feature Importance of each column

3.2 Data Profiling, Cleaning and Preprocessing

Before model training, data profiling was conducted to assess data completeness, consistency, and validity. The following preprocessing steps were performed:

- **Handling Missing Values:**
Missing values in employee skills and availability fields were handled using appropriate imputation strategies such as mean, median, or default category assignment.
- **Data Normalisation:**
Continuous numerical features such as performance scores and availability percentages were normalised to ensure a uniform scale across features.
- **Categorical Encoding:**
Categorical attributes such as department, skill category, and task type were encoded using label encoding or one-hot encoding techniques.
- **Outlier Treatment:**
Extreme values detected during profiling were either capped or removed to prevent skewed model learning.

All preprocessing steps were implemented as reusable pipelines to ensure consistency during training and inference.

3.3 Feature Extraction and Feature Selection

Feature engineering played a crucial role in enhancing model performance. Derived features were created to better represent the relationship between employees and tasks. **Key engineered features include:**

- **Skill Match Score:** Measures the overlap between employee skills and task requirements
- **Availability Score:** Indicates employee workload and current availability
- **Historical Performance Score:** Aggregates past task completion success and quality
- **Task Complexity Score:** Quantifies task difficulty based on required effort and expertise

Feature selection techniques such as correlation analysis and feature importance evaluation were applied to remove redundant or low-impact features. This helped reduce model complexity while maintaining predictive power.




Figure 3.2: The feature correlation heatmap

3.4 Model Design Approaches

The task allocation problem was formulated as a regression problem, where the objective is to predict a continuous suitability score for each employee–task pair. **Three different model design approaches were selected:**

- Linear Regression
- Random Forest Regressor
- XGBoost Regressor

The baseline model helped establish a performance benchmark, while ensemble models were chosen to capture complex, non-linear relationships between features.

3.5 Governance and Data Version Control

To ensure reproducibility and governance, data versioning was implemented using Data Version Control (DVC). Each dataset version was linked to specific Git commits, enabling traceability across experiments. **Governance practices included:**

- Dataset lineage tracking
- Controlled access to sensitive employee data
- Documentation of preprocessing and feature engineering steps

These measures ensure transparency, compliance, and ease of collaboration across the project lifecycle.

Chapter 4

Model Building, Training and Evaluation

4.1 Model Building

Model building involved implementing multiple machine learning algorithms to predict employee–task suitability scores. The following models were developed:

- **Linear Regression:**
Used as a baseline to capture linear relationships between features and suitability score.
- **Random Forest Regressor:**
An ensemble-based model capable of handling non-linear feature interactions and reducing overfitting through bagging.
- **XGBoost Regressor:**
A gradient boosting model designed for high performance and efficient learning from complex data patterns.

All models were implemented using standardized training pipelines to ensure fair comparison.

4.2 Training, Validation and Testing

The dataset was split into training, validation, and testing subsets to evaluate model generalization.

- **Training Set:** Used to fit model parameters
- **Validation Set:** Used for hyperparameter tuning and overfitting detection
- **Test Set:** Used for final unbiased performance evaluation

Training workflows were automated using MLOps pipelines to ensure consistency across experiments. Cross-validation techniques were applied where necessary to improve robustness.

4.3 Evaluation Metrics and Visualization

Model performance was evaluated using regression metrics:

- Mean Absolute Error (MAE)
- Root Mean Squared Error (RMSE)
- R^2 Score (Coefficient of Determination)

These metrics provided insights into both average prediction error and variance explanation. Performance comparisons across models were visualized using graphs generated from MLflow experiment tracking.

Metric/Model	MAE	R ²
Random Forest	0.010	0.9646
Linear Regression	0.012	0.9604
XGBoost	0.0121	0.9611

Table 1. The Evaluation Metrics of the algorithms used

4.4 Experiment Tracking

Experiment tracking was performed using MLflow to log:

- Model parameters
- Evaluation metrics
- Training artifacts
- Model versions

Each experiment run was uniquely identifiable, enabling systematic comparison and selection of the best-performing model. This ensured transparency and reproducibility throughout the development process.




Figure 4.1: MLflow Experiment Tracking

4.5 Interpreting Model Results and Error Analysis

Model interpretability was emphasized to build trust in automated task allocation decisions. Feature importance analysis revealed that:

- Skill match score had the highest influence
- Historical performance and availability followed

Error analysis was conducted to identify scenarios where predictions deviated significantly from actual suitability scores. These insights helped refine feature engineering and model selection strategies.

Model Analysis and Refinement

5.1 Revisiting Model Performance

After the initial development of the Workforce Allocation system, model performance was revisited to ensure reliability and generalisation. Three different machine learning models were trained and evaluated:

- Linear Regression (baseline model)
- Random Forest Regressor
- XGBoost Regressor

Each model was trained using the same feature-engineered dataset containing skill match score, availability, past performance, and task complexity. Performance was evaluated on the training, validation, and test datasets using regression metrics, including the R^2 score, Mean Absolute Error (MAE), and Mean Squared Error (MSE).

The Linear Regression model served as a baseline to understand linear relationships in the data but showed limitations in capturing non-linear dependencies. Random Forest and XGBoost, being ensemble-based models, demonstrated significantly improved predictive performance by modeling complex interactions between features. Comparing training and validation scores helped identify overfitting and underfitting behaviours, ensuring that the selected model performed well on unseen data.

```
✓ Model Evaluation:
R² Score: 0.964
Mean Absolute Error: 0.010
Mean Squared Error: 0.000
Train R²: 0.9674767327780813
Validation R²: 0.9646683725001636
```

Figure 5.1: MAE & R^2 scores for random forest

Metric	Value	Models
MAE	0.01219168082172246	xgboost
R2	0.9604921686627097	xgboost

Figure 5.2: MAE & R^2 scores for XG boost

Metric	Value	Models
MAE	0.012099997929978085	linearregression
R2	0.9611249329860654	linearregression

Figure 5.3: MAE & R^2 scores for linear regression

5.2 Hyperparameter tuning and optimisation

To further improve model robustness, hyperparameter tuning was conducted for both Random Forest and XGBoost models. Parameters such as tree depth, number of estimators, learning rate (for XGBoost), and minimum samples per split were systematically tuned.

Multiple experiments were logged using MLflow, allowing structured comparison of different parameter configurations. The tuning process revealed that constrained tree depth and controlled feature selection significantly reduced overfitting, leading to improved validation performance.

XGBoost showed strong learning capability but required careful tuning to avoid overfitting, whereas Random Forest demonstrated stable performance with comparatively less tuning effort. Linear Regression did not involve extensive tuning due to its simplicity.




Figure 5.4: Model metrics graph from MLflow

5.3 Interpretability

Interpretability was a critical requirement for the Workforce Allocation system to ensure transparency and trust. Feature importance analysis was conducted for ensemble models to identify which attributes most influenced task suitability predictions.

Across Random Forest and XGBoost, skill match score emerged as the most influential feature, followed by past performance and availability. This validated the domain assumption that alignment between employee skills and task requirements plays the most significant role in task assignment.

Linear Regression coefficients further supported these findings by showing positive weight contributions for skill match and performance features. The consistency of feature importance across models strengthened confidence in the dataset design and feature engineering strategy.

5.4 Model versioning and Registry

To manage multiple trained models efficiently, model versioning and registry mechanisms were implemented using MLflow. Each trained model (Linear Regression, Random Forest, XGBoost) was logged along with its parameters, evaluation metrics, and artifacts.

The best-performing model was registered as the production candidate, while alternative models were stored as archived versions. This setup enables easy rollback in case of performance degradation and supports future retraining when new data becomes available.

Model versioning ensures reproducibility and aligns with industry-standard MLOps practices, allowing seamless transition from experimentation to deployment.

5.5 Documentation of Model updates

Comprehensive documentation was maintained throughout the model development lifecycle. Every update including data preprocessing changes, feature engineering steps, algorithm selection rationale, hyperparameter tuning results, and evaluation metrics was recorded.

This documentation ensures transparency, reproducibility, and accountability, making the system easier to maintain and scale. It also enables collaboration across teams and supports auditing and future enhancements.

By integrating documentation with version control and MLflow logs, the project adheres to end-to-end MLOps best practices.

Chapter-6

Model Analysis and refinement

6.1 Overview of CI/CD for ML

Continuous Integration and Continuous Deployment (CI/CD) in Machine Learning extends traditional software CI/CD by incorporating data, models, and experiments into the pipeline. Unlike standard applications, ML systems evolve not only due to code changes but also due to new data, feature updates, and model retraining.

In the Workforce Allocation allocation system, CI/CD ensures that:

- Any update to the dataset, preprocessing logic, or model code automatically triggers training and evaluation
- Model performance is validated before deployment
- Only reliable and well-performing models are promoted to production

CI/CD pipelines reduce manual intervention, minimize errors, and ensure consistent deployment of ML models while maintaining reproducibility and traceability.

6.2 Tools Setup: GitHub Actions

To implement CI/CD for the Workforce Allocation allocation system, popular automation tools such as GitHub Actions, Jenkins, or GitLab CI can be used. These tools automate the execution of workflows triggered by repository events such as code commits or pull requests. GitHub Actions integrates seamlessly with GitHub repositories and is suitable for lightweight ML pipelines.

In this project, the CI/CD tool is configured to:

- Set up the Python environment
- Install dependencies
- Run data preprocessing scripts
- Train and evaluate ML models
- Log metrics and artifacts using MLflow




Figure 7.1: CI automation structure from github

6.3 Automating Model Training and Testing

Automation of model training and testing is a key component of the CI/CD pipeline. Whenever a code change or dataset update occurs, the pipeline automatically executes training scripts for all three models: Linear Regression, Random Forest, and XGBoost.

Each model is evaluated using predefined metrics such as R^2 score, MAE, and MSE. Validation thresholds are enforced to prevent poorly performing models from progressing further in the pipeline. This ensures model quality and guards against performance regression.

Automated testing also includes checks for:

- Data schema consistency
- Missing or invalid feature values
- Overfitting detection using train-validation comparison




Figure 7.2: Metric comparison across CI/CD pipeline runs

6.4 Building and Running a Deployment Pipeline

Once a model passes evaluation checks, it is packaged and prepared for deployment. The deployment pipeline includes steps such as model serialisation, containerization, and environment configuration.

Docker is used to containerise the trained model along with its dependencies, ensuring consistency across environments. The containerised model can then be deployed to local servers or cloud platforms. Model metadata, version details, and performance metrics are tracked using MLflow.

This pipeline ensures that the deployed model is reproducible, scalable, and easily replaceable in case a newer version outperforms the current one.

6.5 Continuous Integration Hands-On Exercise

As part of the hands-on implementation, a CI pipeline was created to automatically train and evaluate the Workforce Allocation allocation models upon each code update. The pipeline performs the following steps:

1. Fetches the latest code from the repository
2. Installs required dependencies
3. Executes preprocessing and feature engineering scripts
4. Trains Linear Regression, Random Forest, and XGBoost models
5. Logs metrics and artifacts using MLflow
6. Validates performance thresholds before deployment

This exercise demonstrated how CI/CD principles can be applied effectively to real-world ML systems, ensuring faster iteration cycles, higher reliability, and improved maintainability.

Chapter-7

Monitoring and tracking the model lifecycle

7.1. Model Monitoring Concepts and KPIs

The model monitoring is critical to ensure that the automated assignment of tasks to employees remains optimal. Unlike static software, machine learning models degrade over time as the environment (employee skills, project requirements) changes. We focus on two distinct categories of Key Performance Indicators (KPIs):

7.1.1. Model Quality Metrics

These measure how accurately the model predicts the **suitability_score** (regression task).

- Coefficient of Determination (R^2 Score): Measures the proportion of variance in the dependent variable (suitability) predictable from the independent variables (skills, availability).
- Root Mean Squared Error (RMSE): This penalizes larger errors more heavily, ensuring that the model does not make catastrophic misallocations.
- Prediction Distribution: Monitoring the statistical distribution of predicted scores ensures the model is not outputting only extreme values (e.g., always predicting 100% or 0% suitability).

7.1.2. Operational Metrics

- Inference Latency: The time taken to generate a task recommendation. High latency can disrupt the user experience in the allocation dashboard.
- Data Quality: Identifying missing values or type mismatches in incoming feature vectors (e.g., a missing `skill_id`) before inference occurs.

7.2. Detecting Concept Drift and Data Drift

Evidently is integrated directly into the Python inference pipeline. It generates JSON and HTML reports that visualize distribution changes. Specifically, the `DataDriftPreset` and `TargetDriftPreset` are utilized to calculate statistical distances between the reference dataset (training data) and the current production batch. Drift detection is the core mechanism for deciding when the model is no longer valid. In the Workforce Allocation Allocation System, we distinguish between two specific types:

7.2.1. Data Drift (Covariate Shift)

This occurs when the statistical properties of the input features change.

- Scenario: A sudden influx of tasks requiring a new technology (e.g., "Rust") that was not present in the training data, or a shift in employee availability patterns due to holiday seasons.
- Detection Method: Evidently AI utilizes the Kolmogorov-Smirnov (K-S) test for numerical features (e.g., experience_years) and the Chi-Square test for categorical features (e.g., department_id) to detect significant divergence.

7.2.2. Concept Drift

This occurs when the relationship between features and the target variable changes

- Scenario: The organization changes its definition of "suitability." Previously, technical skill was the highest weight; now, availability is the priority. The old model will continue predicting high scores for unavailable experts, which is now incorrect.
- Detection Method: We monitor the correlation between predicted scores and actual feedback ratings (ground truth) collected post-task completion.

Column	Type	Reference Distribution	Current Distribution	Data Drift	Stat Test	Drift Score
> skill_match_score	num			Not Detected	Wasserstein distance (normed)	0.007453
> availability	num			Not Detected	Wasserstein distance (normed)	0.006835
> past_performance	num			Not Detected	Wasserstein distance (normed)	0.006177
> suitability_score	num			Not Detected	Wasserstein distance (normed)	0.003369
> complexity_level	num			Not Detected	Jensen-Shannon distance	0.001892
> complexity_inv	num			Not Detected	Jensen-Shannon distance	0.001892

Figure 8.1: Dataset drift visualization using EvidentlyAI

7.3. Model Retraining and Redeployment

To maintain performance, the system implements an automated Continuous Training (CT) pipeline orchestrated by Prefect.

- Trigger: Retraining is triggered either on a schedule (e.g., monthly) or conditionally when Evidently AI detects drift exceeding the threshold (e.g., drift detected in >30% of features).
- Data Versioning: DVC pulls the latest validated dataset, ensuring the training data is traceable and reproducible.

- Training Pipeline: Prefect executes the training script, which includes preprocessing, feature engineering, and model fitting (XGBoost/Random Forest).
- Evaluation: The new candidate model is evaluated against a hold-out validation set.
- Registration: If the new model performs better than the current production model, it is registered in the MLflow Model Registry under the "Staging" stage.
- Deployment: A CI/CD pipeline builds a new Docker container with the updated model artifact and deploys it to the inference server.

7.4. Logging, Alerting, and Dashboard Setup

Effective observability requires granular logging and proactive alerting.

- MLflow Tracking: Every training run logs:
 - Parameters: `max_depth`, `n_estimators`, `learning_rate`.
 - Metrics: Training and Validation Loss, MAE, R^2 .
 - Artifacts: The serialized model (pickle/joblib), feature importance plots, and dependency requirements.
- Alerting Rules:

Alerts are configured to notify the MLOps team via email or Slack if:

 - The percentage of drifting features exceeds 50%.
 - Inference latency exceeds 500ms for the 95th percentile (P95).
 - The prediction API returns >5% error rate within a 5-minute window.
- Unified Dashboard: A Streamlit-based interface integrates Evidently's HTML reports, allowing non-technical stakeholders to visually inspect why the model might be behaving differently.




Figure 8.2: Logging using Prefect

Chapter-8

Performance Evaluation and Governance


8.1. Post-Deployment Performance Review

After deployment, the Intelligent Workforce Allocation System undergoes a structured post-deployment performance review using production-like or live operational data. The objective of this review is to ensure that the deployed model continues to perform as expected in real-world conditions.

All relevant evaluation metrics such as MAE, R^2 , prediction stability, and task allocation accuracy are logged and tracked using MLflow. Multiple model versions are compared side-by-side to assess improvements or regressions over time. Based on these comparisons, the best-performing and most stable model version is selected for continued use in production. This review process ensures:


- Accurate task-to-employee matching, improving task success rates
- Stable prediction performance under changing workloads and data distributions
- Reduced task misallocation, minimizing employee overload and inefficiencies

By systematically reviewing post-deployment metrics, the system avoids silent degradation and maintains consistent operational quality.



The screenshot displays the 'Intelligent Workforce Allocation System' interface. At the top, it says 'Compare predictions from multiple ML models'. Below this is the 'Input Details' section, which includes three input fields: 'Employee Skills (comma-separated)' with the value 'python, django, sql', 'Required Task Skills (comma-separated)' with the value 'python, flask, sql', and 'Task Complexity' set to 'Easy'. At the bottom, there are two horizontal progress bars: 'Availability' with a red dot at 0.80 and 'Past Performance' with a red dot at 0.90.

Figure 8.1: Input details and Model prediction



The screenshot displays the MLflow Model Monitoring & Observability interface. It features a table titled 'MLflow Model Metrics' with columns for Run Name, MAE, R2, MSE, and MAPE. The table lists four runs: 'dazzling_hound_470', 'Random_Forest', 'XGBoost_Regressor', and 'Linear_Regression'. The 'dazzling_hound_470' run has 'None' values for all metrics. The other three runs have numerical values for all metrics.

Run Name	MAE	R2	MSE	MAPE
0 dazzling_hound_470	None	None	None	None
1 Random_Forest	0.0126	0.956	0.0003	0.0294
2 XGBoost_Regressor	0.012	0.9619	0.0002	0.028
3 Linear_Regression	0.0119	0.9625	0.0002	0.0278

Figure 8.2: Model monitoring integrating MLflow

8.2. Continuous Improvement via Feedback Loops

To ensure long-term effectiveness, the system incorporates **continuous feedback loops**. After task completion, **task outcomes**, **completion time**, and **employee performance indicators** are collected and stored. This real-world feedback is then integrated into future training cycles, allowing the model to learn from past decisions. The continuous improvement process includes:

- **Updating datasets** with newly collected task and performance data
- **Re-running automated Prefect pipelines** for data processing, training, and evaluation
- **Comparing newly trained models with previous versions using MLflow**, ensuring only improved models are promoted

Through this iterative loop, the system progressively improves task allocation quality, adapts to changing workforce dynamics, and remains aligned with organizational goals.

8.3. Governance and Ethical AI Practices

Strong governance practices are implemented to ensure that the Intelligent Workforce Allocation System operates in a **fair, transparent, and responsible manner**. Governance mechanisms focus on both technical reliability and ethical considerations. Key governance objectives include:

- **Fair task allocation across employees**, preventing systematic overloading or underutilization
- **Transparent model decision-making**, enabling explainability of task assignments
- **Secure handling of employee and task data**, ensuring privacy and access control
- **Prevention of biased or discriminatory recommendations**

Ethical AI practices are enforced by carefully monitoring **feature importance**, validating that no sensitive or proxy attributes disproportionately influence predictions. Regular reviews are conducted to ensure compliance with fairness and responsible AI principles.

8.4. Documentation: Model Cards and Data Cards

To improve transparency and audit readiness, the project maintains comprehensive documentation in the form of **Model Cards** and **Data Cards**.

Model Cards

- The **purpose and scope** of the Intelligent Workforce Allocation model
- The **algorithms used** (e.g., Linear Regression, XGBoost)
- **Performance metrics** across different evaluation stages
- **Known limitations and assumptions**

Data Cards

- Dataset composition (employee skills, task attributes, outcomes)
- Data sources and **preprocessing steps**
- Feature engineering logic
- Potential **data biases or coverage gaps**

Together, these artifacts enhance **reproducibility, transparency, and stakeholder trust**.

8.5. Audit and Review Checklist

A structured audit and review process is followed to ensure long-term system reliability and compliance. The audit checklist includes:

- Verification of **MLflow experiment logs and model lineage**
- Review of **Evidently AI drift and data quality reports**
- Validation of **automated retraining pipelines** managed by Prefect
- Confirmation of **Dockerized deployment consistency** across environments
- Review of **documentation completeness**, including Model Cards and Data Cards

Regular audits help detect operational issues early, maintain compliance with governance standards, and ensure that the system continues to deliver reliable and ethical task allocation outcomes over time.

Appendix

A.1 List Of MLOps Tools, Its Features, Pros And Cons

1. MLflow: Experiment tracking, model versioning, and artifact management. Key Features:

- Logs parameters, metrics, and model artifacts
- Enables reproducibility of model training
- Supports model comparison and evaluation

Pros

- Lightweight and easy to integrate
- Strong support for experiment reproducibility
- Clear traceability between models and metrics

Cons

- Model registry features require additional setup
- Not designed for real-time monitoring

2. Prefect: Workflow orchestration for monitoring, drift detection, and retraining. Key Features:

- Task and flow-based pipeline orchestration
- Scheduling and retry mechanisms
- Observability of ML workflows
- Trigger-based pipeline execution

Pros

- Highly flexible and Python-native
- Easy to debug and monitor workflows
- Strong support for conditional execution
- Ideal for complex ML pipelines

Cons

- Needs external tools for visualization and metrics
- Adds orchestration overhead for small pipelines

3. Docker: Containerization of inference service and frontend. Key Features:

- Environment consistency across development and production
- Isolated runtime for inference and frontend
- Image-based deployment
- Dependency encapsulation

Pros

- Ensures reproducible deployments
- Simplifies environment setup
- Reduces system compatibility issues
- Industry-standard deployment practice

Cons

- Requires container management expertise
- Debugging containers can be complex
- Slight performance overhead

4. EvidentlyAI: Data drift and model behavior analysis. Key Features:

- Data drift detection
- Prediction distribution analysis
- Feature-level statistical comparison
- Automated report generation

Pros

- ML-specific monitoring capabilities
- Easy integration with logged data
- Clear, interpretable reports
- Open-source friendly

Cons

- Not real-time by default
- Computationally expensive for large datasets

5. Git & GitHub: Source code version control and collaboration. Key Features:

- Code versioning
- Branching and merging
- CI/CD pipeline triggers
- Collaboration support

Pros

- Industry-standard version control
- Enables CI/CD automation
- Strong collaboration features

Cons

- No native model or data versioning
- Requires discipline in workflow management
- Limited visibility into ML artifacts

6. DVC (Data Version Control): Dataset and artefact versioning. Key Features:

- Data versioning alongside code
- Lightweight tracking of large datasets
- Pipeline reproducibility
- Storage abstraction

Pros

- Enables reproducible ML experiments
- Integrates well with Git
- Scalable for large datasets

Cons

- Initial setup complexity
- Requires external storage backend

A.2 Tool Installation Guides

1. MLFlow

- Create and activate a virtual environment
- Install MLflow using pip
- Verify installation by running the MLflow CLI and checking the version

2. DVC

- Install DVC using pip install dvc
- Initialize DVC inside the Git repository using dvc init
- Configure a local or remote storage backend (optional)

3. PREFECT

- Install Prefect using pip install prefect
- Verify installation by running prefect version
- Initialize Prefect locally (optional) for UI and flow tracking.

4. EVIDENTLY AI

- Install EvidentlyAI using pip install evidently
- Import Evidently modules within drift detection scripts

5. DOCKER

- Download Docker Desktop (Windows/macOS) or Docker Engine (Linux)
- Enable Docker daemon
- Verify installation using docker --version

6. GIT & GITHUB


- Install Git using OS-specific installers
- Configure Git with username and email
- Create a GitHub repository and link it to the local project

A.3 References

- [1] A. Sharma and P. Kumar, "MLOps: Changes and Tools – Adapting Machine Learning Workflows for Enhanced Efficiency and Scalability," *International Journal of Engineering Research & Technology (IJERT)*, vol. 14, no. 11, pp. 1–6, Nov. 2025.
- [2] K. Patel and S. J. Rao, "Data Versioning for Continuous AI Deployment," *International Journal of Advanced Computer Science and Applications*, vol. 16, no. 7, pp. 210–218, Jul. 2025.
- [3] J. Nilsson, "Real-Time Machine Learning Model Monitoring for Banking Fraud Detection," M.Sc. Thesis, Dept. Info. Tech., Uppsala Univ., Uppsala, Sweden, 2025.
- [4] R. Müller, M. Abdelaal, and D. Stjelja, "Open-Source Drift Detection Tools in Action: Insights from Two Use Cases," in *arXiv preprint arXiv:2404.18673*, Apr. 2024.
- [5] D. Al-Fraihat, Y. Sharrab, A. Al-Ghuwairi, H. Alzabut, M. Beshara, and A. Algarni, "Utilizing machine learning algorithms for task allocation in distributed agile software development," *Heliyon*, vol. 10, no. 21, e39926, 2024.
- [6] N. Grigoryan, "Comparative Analysis of Open-Source ML Pipeline Orchestration Platforms," B.Sc. Thesis, Dept. Computer Sci., Univ. of Trier, Germany, 2024.
- [7] S. Sharma and R. Gupta, "Intelligent Task Management System," *International Journal of Creative Research Thoughts (IJCRT)*, vol. 11, no. 12, pp. 456–462, Dec. 2023.
- [8] V. Nakra, A. Dave, B. Devaguptapu, P. K. Chenchala, and A. Agarwal, "Enhancing Software Project Management and Task Allocation with AI and Machine Learning," *International Journal on Recent and Innovation Trends in Computing and Communication*, vol. 11, no. 11, pp. 1171–1178, Nov. 2023.
- [9] S. Gupta and R. S. Rao, "LeSS agile projects: a machine learning driven empirical model to predict the human resource allocation," *International Journal of Information Technology*, vol. 16, no. 5, pp. 1–10, Oct. 2023.
- [10] A. K. S. and P. R., "Efficient Workforce Planning with a Labour Allocation System," *International Journal for Research in Applied Science and Engineering Technology*, vol. 11, no. 6, pp. 1–5, Jun. 2023.
- [11] H. Grover and S. Sharma, "Machine Learning Algorithms and Predictive Task Allocation in Software Project Management," *International Journal of Open Publication and Exploration (IJOPE)*, vol. 11, no. 1, p. 34, Jan. 2023.
- [12] V. T. P. N. and K. C., "A Novel Multi-Objective Evolutionary Algorithm to Address Turnover in Software Project Scheduling," in *2023 IEEE International Conference on Computing, Communication and Networking Technologies (ICCCNT)*, Delhi, India, 2023, pp. 1–6.

- [13] R. J. Wave, "AI Algorithm for Workforce Scheduling," *International Journal of Engineering Development and Research*, vol. 10, no. 2, pp. 169–175, Jun. 2022.
- [14] A. Sadegheih, A. Nazari, and R. Tehrani, "Project scheduling based on resource leveling using fuzzy ranking and genetic algorithm: Examination and analysis," *International Journal of Health Sciences*, vol. 6, no. S8, pp. 409–432, May 2022.
- [15] D. Mourtzis, V. Siatras, and J. Angelopoulos, "An intelligent model for workforce allocation taking into consideration the operator skills," *Procedia CIRP*, vol. 97, pp. 196–201, 2021.
- [16] P. K. and S. R., "An Analysis of Deep Neural Network for Recommending Developers to Fix Reported Bugs," in *2021 International Conference on Artificial Intelligence and Smart Systems (ICAIS)*, Coimbatore, India, 2021, pp. 1–6.
- [17] S. K. and R. B., "Reinforcement Learning in Dynamic Task Scheduling: A Review," *International Journal of Computer Applications*, vol. 176, no. 5, pp. 15–21, Oct. 2020.
- [18] M. A. and S. H., "Fuzzy Logic-Based Algorithm Resource Scheduling For Improving The Reliability Of Cloud Computing," *International Journal of Scientific & Technology Research*, vol. 9, no. 3, pp. 120–125, Mar. 2020.
- [19] M. Zaharia, A. Chen, A. Davidson, A. Ghodsi, S. A. Hong, et al., "Accelerating the Machine Learning Lifecycle with MLflow," in *IEEE Data Engineering Bulletin*, vol. 41, no. 4, pp. 39–45, Jun. 2018.
- [20] F. Roque et al., "Task Allocation for Crowdsourcing Using AI Planning," in *2016 IEEE/ACM 3rd International Workshop on CrowdSourcing in Software Engineering (CSI-SE)*, Austin, TX, USA, 2016, pp. 1–7.

A.4 Folder Structure



A.5 Further Reading

1. DVC – <https://dvc.org>
2. MLflow – <https://mlflow.org>
3. Prefect – <https://www.prefect.io>
4. GitHub Actions – <https://docs.github.com/actions>
5. Docker – <https://docs.docker.com>
6. Evidently AI – <https://www.evidentlyai.com>

Dataset Link : <http://kaggle.com/datasets/umerfarooq09/skill-based-task-assignment>

Github Link : https://github.com/loki-ram/MLops_EL